

In [1]:

```
# This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

# Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = "Colab Notebooks/assignment1"
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd drive/My Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My Drive/$FOLDERNAME
```

```
Mounted at /content/drive
/content/drive/My Drive/Colab Notebooks/assignment1/cs231n/datasets
/content/drive/My Drive/Colab Notebooks/assignment1
```

Softmax exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

This exercise is analogous to the SVM exercise. You will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

In [2]:

```
import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

In [3]:

```
def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000, num_dev=
```

```
"""
```

```
Load the CIFAR-10 dataset from disk and perform preprocessing to prepare  
it for the linear classifier. These are the same steps as we used for the  
SVM, but condensed to a single function.
```

```
"""
```

```
# Load the raw CIFAR-10 data
```

```
cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'
```

```
# Cleaning up variables to prevent loading data multiple times (which may cause me  
try:
```

```
    del X_train, y_train
```

```
    del X_test, y_test
```

```
    print('Clear previously loaded data.')
```

```
except:
```

```
    pass
```

```
X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)
```

```
# subsample the data
```

```
mask = list(range(num_training, num_training + num_validation))
```

```
X_val = X_train[mask]
```

```
y_val = y_train[mask]
```

```
mask = list(range(num_training))
```

```
X_train = X_train[mask]
```

```
y_train = y_train[mask]
```

```
mask = list(range(num_test))
```

```
X_test = X_test[mask]
```

```
y_test = y_test[mask]
```

```
mask = np.random.choice(num_training, num_dev, replace=False)
```

```
X_dev = X_train[mask]
```

```
y_dev = y_train[mask]
```

```
# Preprocessing: reshape the image data into rows
```

```
X_train = np.reshape(X_train, (X_train.shape[0], -1))
```

```
X_val = np.reshape(X_val, (X_val.shape[0], -1))
```

```
X_test = np.reshape(X_test, (X_test.shape[0], -1))
```

```
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))
```

```
# Normalize the data: subtract the mean image
```

```
mean_image = np.mean(X_train, axis = 0)
```

```
X_train -= mean_image
```

```
X_val -= mean_image
```

```
X_test -= mean_image
```

```
X_dev -= mean_image
```

```
# add bias dimension and transform into columns
```

```
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
```

```
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
```

```
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
```

```
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])
```

```
return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev
```

```
# Invoke the above function to get our data.
```

```
X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev = get_CIFAR10_data()
```

```
print('Train data shape: ', X_train.shape)
```

```
print('Train labels shape: ', y_train.shape)
```

```
print('Validation data shape: ', X_val.shape)
```

```
print('Validation labels shape: ', y_val.shape)
```

```
print('Test data shape: ', X_test.shape)
```

```
print('Test labels shape: ', y_test.shape)
```

```
print('dev data shape: ', X_dev.shape)
```

```
print('dev labels shape: ', y_dev.shape)
```

```
Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)
```

Softmax Classifier

Your code for this section will all be written inside `cs231n/classifiers/softmax.py`.

```
In [5]: # First implement the naive softmax loss function with nested loops.
# Open the file cs231n/classifiers/softmax.py and implement the
# softmax_loss_naive function.

from cs231n.classifiers.softmax import softmax_loss_naive
import time

# Generate a random softmax weight matrix and use it to compute the loss.
W = np.random.randn(3073, 10) * 0.0001
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As a rough sanity check, our loss should be something close to -log(0.1).
print('loss: %f' % loss)
print('sanity check: %f' % (-np.log(0.1)))
```

```
loss: 2.334964
sanity check: 2.302585
```

Inline Question 1

Why do we expect our loss to be close to $-\log(0.1)$? Explain briefly.**

Your Answer : 우리가 작성한 코드를 보면 *label class*의 수는 10개이고, 가중치를 *random* 함수를 사용해 무작위로 설정하였으므로, 평균적으로 1/10의 정확도를 가지게 된다. 그러므로 수식에 따라 $-\log(1/10) = -\log(0.1)$ 에 가까운 *loss*를 갖게 된다.

```
In [6]: # Complete the implementation of softmax_loss_naive and implement a (naive)
# version of the gradient that uses nested loops.
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As we did for the SVM, use numeric gradient checking as a debugging tool.
# The numeric gradient should be close to the analytic gradient.
from cs231n.gradient_check import grad_check_sparse
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)

# similar to SVM case, do another gradient check with regularization
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

```
numerical: 0.663658 analytic: 0.663658, relative error: 3.057256e-08
numerical: 0.863740 analytic: 0.863740, relative error: 6.476761e-08
numerical: -0.140858 analytic: -0.140858, relative error: 1.135148e-07
numerical: -0.002281 analytic: -0.002281, relative error: 2.127957e-06
numerical: -0.731742 analytic: -0.731742, relative error: 7.638961e-09
numerical: -0.131487 analytic: -0.131487, relative error: 3.927650e-08
numerical: 0.868624 analytic: 0.868624, relative error: 4.992314e-09
numerical: 1.656665 analytic: 1.656665, relative error: 1.203253e-08
```

```

numerical: -2.508937 analytic: -2.508937, relative error: 9.879609e-09
numerical: -2.908926 analytic: -2.908926, relative error: 2.805085e-08
numerical: 1.088320 analytic: 1.088320, relative error: 8.855573e-09
numerical: 0.478719 analytic: 0.478719, relative error: 1.649932e-08
numerical: -0.449663 analytic: -0.449663, relative error: 9.521726e-08
numerical: 0.999332 analytic: 0.999332, relative error: 9.147471e-09
numerical: 0.307334 analytic: 0.307334, relative error: 1.249076e-08
numerical: 0.671166 analytic: 0.671166, relative error: 4.660121e-08
numerical: 1.293835 analytic: 1.293835, relative error: 6.446288e-09
numerical: 2.931522 analytic: 2.931522, relative error: 1.318756e-08
numerical: 3.886253 analytic: 3.886253, relative error: 3.958368e-09
numerical: 3.456118 analytic: 3.456118, relative error: 1.277052e-09

```

In [7]:

```

# Now that we have a naive implementation of the softmax loss function and its gradient
# implement a vectorized version in softmax_loss_vectorized.
# The two versions should compute the same results, but the vectorized version should
# much faster.
tic = time.time()
loss_naive, grad_naive = softmax_loss_naive(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('naive loss: %e computed in %fs' % (loss_naive, toc - tic))

from cs231n.classifiers.softmax import softmax_loss_vectorized
tic = time.time()
loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

# As we did for the SVM, we use the Frobenius norm to compare the two versions
# of the gradient.
grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
print('Loss difference: %f' % np.abs(loss_naive - loss_vectorized))
print('Gradient difference: %f' % grad_difference)

```

```

naive loss: 2.334964e+00 computed in 0.148034s
vectorized loss: 2.334964e+00 computed in 0.018492s
Loss difference: 0.000000
Gradient difference: 0.000000

```

In [18]:

```

# Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of over 0.35 on the validation set.

from cs231n.classifiers import Softmax
results = {}
best_val = -1
best_softmax = None

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength.
# This should be identical to the validation that you did for the SVM; save
# the best trained softmax classifier in best_softmax.
#####

# Provided as a reference. You may or may not want to change these hyperparameters
learning_rates = [1e-7, 5e-7]
regularization_strengths = [2.5e4, 5e4]

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

from itertools import product
grid_search = list(product(*[learning_rates, regularization_strengths]))

```

```

for lr, rg in grid_search :
    soft = Softmax()

    soft.train(X_train, y_train, learning_rate = lr, reg = rg,
               num_iters = 1000)

    y_train_predict = soft.predict(X_train)
    learning_accuracy = np.mean(y_train_predict == y_train)

    y_validation_predict = soft.predict(X_val)
    validation_accuracy = np.mean(y_validation_predict == y_val)

    results[(lr,rg)] = (learning_accuracy, validation_accuracy)

    if best_val < validation_accuracy :
        best_val = validation_accuracy
        best_softmax = soft

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' % best_val)

```

```

lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.327000 val accuracy: 0.340000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.309673 val accuracy: 0.320000
lr 5.000000e-07 reg 2.500000e+04 train accuracy: 0.332939 val accuracy: 0.351000
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.293857 val accuracy: 0.307000
best validation accuracy achieved during cross-validation: 0.351000

```

In [19]:

```

# evaluate on test set
# Evaluate the best softmax on test set
y_test_pred = best_softmax.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))

```

```
softmax on raw pixels final test set accuracy: 0.351000
```

Inline Question 2 - True or False

Suppose the overall training loss is defined as the sum of the per-datapoint loss over all training examples. It is possible to add a new datapoint to a training set that would leave the SVM loss unchanged, but this is not the case with the Softmax classifier loss.

Your Answer : True

Your Explanation : 일반적으로, SVM은 robust하며 어느정도 이상의 성능을 갖추게 되면 loss 가 잘 변하지 않는다. 하지만, Softmax는 정확도가 1에 근접할 때까지 계속 성능을 계속하는 것이 특징이므로 True이다.

In [20]:

```

# Visualize the learned weights for each class
w = best_softmax.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)

w_min, w_max = np.min(w), np.max(w)

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for i in range(10):

```

```
plt.subplot(2, 5, i + 1)

# Rescale the weights to be between 0 and 255
wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
plt.imshow(wimg.astype('uint8'))
plt.axis('off')
plt.title(classes[i])
```

